Multi-Stage Ensemble and Feature Engineering for MOOC Dropout Prediction

Jeong-Yoon Lee
Conversion Logic

12300 Wilshire Blvd. Los Angeles, CA 90025, USA

jeong@conversionlogic.com

Andreas Toescher
Opera Solutions

Hauptplatz 12, 8580 Koeflach, Austria

andreas.toescher@commendo.at

Michael Jahrer
Opera Solutions

Hauptplatz 12, 8580 Koeflach, Austria

michael.jahrer@commendo.at

Kohei Ozaki
AIG Japan Holdings KK

Kamiyacho MT Bldg 18F, 4-3-20 Toranomon, Minato-ku, Tokyo 105-0001, Japan

ozaki.kohei@aig.co.jp

Mert Bay
Conversion Logic

12300 Wilshire Blvd. Los Angeles, CA 90025, USA

mert@conversionlogic.com

Peng Yan
NetEase Youdao

Chuangye Building, Tsinghua Science Park, Beijing, 100084, China

yanpeng@rd.netease.com

Song Chen
American International Group

175 Water Street, New York, NY 10038, USA

song.chen@aig.com

Tam T. Nguyen
Institute for Infocomm Research, A*STAR

1 Fusionopolis Way, #21-01 Connexis, Singapore 138632

nguyentt@i2r.a-star.edu.sg

Xiaocong Zhou
Tsinghua University

Haidian District, Beijing, 100084, China

infinitezxc@gmail.com

ABSTRACT

This paper describes the winning solution of KDD Cup 2015. The competition aims to predict dropouts in Massive Open Online Courses (MOOCs). Our approach begins with feature engineering to extract predictive features from activity logs of students and meta data. Then, we train sixty three individual classifiers with different subsets of features and seven algorithms. Lastly, we blend predictions of individual classifiers with the multi-stage ensemble framework. Our solution achieves AUC scores of 0.90918 and 0.90744 on the public and private leaderboards respectively.

Categories and Subject Descriptors

I.5.4 [**Pattern Recognition**]: Application

General Terms

Application

Keywords

KDD Cup, Feature Engineering, Ensemble Learning

1. INTRODUCTION

The task of KDD Cup 2015 is to predict the likelihood of dropout for students on XuetangX, one of the largest Massive Open Online Course (MOOC) platforms in China.

Activity logs of 200,906 enrollments from 112,448 students across 39 courses are provided. Each activity is described by 6 fields of the username, course ID, timestamp, source, event, and object. For each object, 3 additional fields of the category, children, and start date are provided. The training set consists of 8,157,278 logs from 120,543 enrollments with the target variable indicating if a student dropped out. The test set consists of 5,387,848 logs from 80,363 enrollments. The full description of the data sets is available in [**?**]

Our final solution is a joint work from 9 data scientists, distributed around the world. The pipeline from raw data to final solution is as follows:

- Hand crafted feature engineering (most of hard work)
- Automatic feature design (autoencoder)
- Individual models (gbm, nn, factor model,..)
- Stage-I ensemble (blends individual models)
- Stage-II ensemble (blends stage-I ensemble models)
- Stage-III ensemble (blends stage-II ensemble models)

2. FEATURE ENGINEERING

All features.

2.1 Feature SK

Features generated by Song and Kohei can be classified as follows:

- Enrollment-based features (No.1-8)
- Username-based features (No.9-18)
- Username-based features for each courses (No.19-25)

- Features based on 10 days after the end date of course (No.26-35)
- Features based on 1 day after the end date of a course (No.36-45)
- Day-level features (No.46)
- Day-level features using target variables (No.47-58)

Full list of features generated by Song and Kohei are described in Table 1. (just listing them for now. TBD in detail).

2.2 Feature RW

Peng and Xiaocong feature.

2.3 Feature TN

Tam feature.

2.4 Feature MJ

Features generated by Michael Jahrer are in sparse format:

- uID (0-112,447)
- cID (112,448-112,486)
- uIDcnt (112,487-112,487)
- eIDcnt (112,488-112,488)
- eID $\rightarrow$ sID (112,489-112,490)
- eID $\rightarrow$ evID (11,2491-112,497)
- eID $\rightarrow$ oIDCnt (112,498-139,443)
- eID $\rightarrow$ tIDCnt (139,444-139,635)
- uID: floor(log(dateSpan^2+1)) (139,636-140,635)
- uID $\rightarrow$ log(time diff to obj start+1) (140,636-140,636)
- eID $\rightarrow$ dateVec diff stats (140,637-140,649)

2.5 Feature MB

Mert feature.

2.6 Feature JL

Features generated by Jeong-Yoon Lee are as follows:

- User ID (20,113) - One-hot-encoded username. Usernames appearing less than 100 times in training log data are grouped together as one user ID.
- Course ID (39) - One-hot-encoded course_id.
- Source Event (10) - One-hot-encoded combination of source and event.
- Object ID (3,554) - One-hot-encoded object. Objects appearing less than 100 times in training log data are grouped together as one object ID.
- Count (1) - Number of log entries for an enrollment_id.
- Object Category (6) - Number of log entries with an object category for an enrollment_id.
- Number of Children Objects (7) - One-hot-encoded total number of object's children for an enrollment_id.
- Object Timespan (10) - One-hot-encoded timespan in days between object's start date and last day of the class

- Day of Class (30) - One-hot-encoded day of the class
- Week of Class (4) - One-hot-encoded week of the class
- End Month of Class (7) - One-hot-encoded end month of the class
- Object Started in Dropout Period (2) - Binary variable that is 1 if object started after but before 10 days after last day of the class and 0 otherwise.

3. MODEL VALIDATION

We use stratified 5-fold cross validation to estimate the performance of single and ensemble models: Training data are split into five folds while the sample size and dropout rate are preserved across folds.

For validation, each of single and ensemble models is trained five times. Each time, one fold is held out and the remaining four folds are used for training. Then, predictions for the hold-out folds are combined and form the model's CV prediction. CV predictions are used in AUC score calculation and/or as inputs in ensemble model training.

For test, each of single and ensemble models is retrained with whole training data. Then predictions for test data are used for submission and/or as inputs in ensemble model prediction.

4. SINGLE MODELS

something like this.

4.1 Learning Algorithms

- Logistic Regression (LR)
- Kernel Ridge Regression (KRR)
- Factorization Machine/Field-aware Factorization Machine (FM/FFM)
- Neural Networks (NN)
- Extreme Trees (ET)
- Gradient Boosting Decision Trees (GBDT)

4.2 Single Models

- Model 1: RandomForest(R). Dataset: X
- Model 2: Logistic Regression(scikit). Dataset: Log(X+1)
- Model 3: Extra Trees Classifier(scikit). Dataset: Log(X+1) (but could be raw)
- Model 4: KNeighborsClassifier(scikit). Dataset: Scale(Log(X+1))
- Model 5: libfm. Dataset: Sparse(X). Each feature value is a unique level.

5. MULTI-STAGE ENSEMBLE

Stratified 5-fold cross validation (CV). We used xgboost [?], neural nets and linear regression for stage-II ensembling.

5.1 Stage-I Ensemble

We trained 20 stage-I ensemble classifiers with different subsets of CV predictions of 63 individual classifiers.

5.2 Stage-II Ensemble

We trained 2 stage-II ensemble classifiers with different subsets of CV predictions of 20 stage-I ensemble classifiers.

5.3 Stage-III Ensemble

We trained a stage-III ensemble classifier with CV predictions of 5 classifiers: 2 stage-II ensemble, 2 stage-I ensemble, and 1 individual classifiers.

id	name	type	5CV	linear weight
1	trn.final.90788	Stage-I	0.907878	1.96267
2	esb58v5+magic.dae+nn.validCV.0.907567	Stage-I	0.907567	0.787138
3	et_esb58v5_rank.val.0.906207	Stage-I	0.906207	0.458095
4	lr_forward_0.01_esb.esb15v3.val.yht	Stage-II	0.907968	1.61461
5	xgb_rf.ko_new_feat.txt.valCV.0.906721	Single	0.906721	1.1703

A linear combination of the 5 models from table 5.3 results in train AUC=0.908072 and accuracy=0.887334. Which leads to 0.90910 public leaderboard score. By adding 39 courseID correction factors train AUC=0.908194 and public score improved to 0.90918.

6. CONCLUSIONS

Our final AUC score of 0.90918 results from a complex pipeline from raw data to final score. Every part of that pipe needs to be (sub-)optimal implemented by our team to get the best score at the end. The first part "feature design" is the most important one and needs expertise, experience and of course a bit luck to capture all signals in the data.

7. ACKNOWLEDGEMENTS

No.	Description
1	Course_id encoded by 1-of-N coding
2	Number of requests by an enrollment_id
3	Number of unique object by an enrollment_id
4	Number of unique problem object of event by an enrollment_id
5	Number of active days by an enrollment_id
6	Number of active hours by an enrollment_id
7	Time of first access in hours by an enrollment_id
8	Time of last access in hours by an enrollment_id
9	Number of enrollments by an username
10	Number of requests by an username
11	Number of unique objects by an username
12	Number of unique problem object of event by an username
13	Number of active days by an username
14	Number of active hours by an username
15	Time of first access in hours by an username
16	Time of last access in hours by an username
17	Time of first problem access in hours by an username
18	Time of last problem access in hours by an username
19	For each course, number of requests by an username
20	For each course, number of unique object by an username
21	For each course, number of unique problem object by an username
22	For each course, number of active days by an username
23	For each course, number of active hours by an username
24	For each course, time of first access in hours
25	For each course, time of last access in hours
26	Number of enrollment_ids during 10 days after the end date of course by an username
27	For each course, number of access logs during 10 days after the end date of course by an username
28	For each course, number of unique objects during 10 days after the end date of course by an username
29	For each course, number of unique problem objects during 10 days after the end date of course by an username
30	For each course, number of active hours during 10 days after the end date of course by an username
31	For each course, difference between first and last access during 10 days after the end date of course by an username
32	For each course, time of first access in hours during 10 days after the end date of course by an username
33	For each course, time of last access in hours during 10 days after the end date of course by an username
34	For each course, time of first access to an problem object in hours during 10 days after the end date of course by an username
35	For each course, time of last access to an problem object in hours during 10 days after the end date of course by an username
36	Number of enrollment_ids during 1 day after the end date of course by an username
37	For each course, number of access logs during 1 day after the end date of course by an username
38	For each course, number of unique objects during 1 day after the end date of course by an username
39	For each course, number of unique problem objects during 1 day after the end date of course by an username
40	For each course, number of active hours during 1 day after the end date of course by an username
41	For each course, difference between first and last access during 1 day after the end date of course by an username
42	For each course, time of first access in hours during 1 day after the end date of course by an username
43	For each course, time of last access in hours during 1 day after the end date of course by an username
44	For each course, time of first access to an problem object in hours during 1 day after the end date of course by an username
45	For each course, time of last access to an problem object in hours during 1 day after the end date of course by an username
46	For each days of the course, which date is provided in date.csv, number of unique active courses by an username
47	For each 10 days after the end date of the course, number of active enrollment_id, which target variables are 1 in the training set, enrolled by
48	For each 10 days after the end date of the course, number of active enrollment_id, which target variables are 0 in the training set, enrolled by
49	For each 10 days after the end date of the course, number of active enrollment_id (in this case, days between last access and the end date of t
50	For each 10 days after the end date of the course, number of active enrollment_id (in this case, days between last access and the end date of t
51	For each 14 days before the end date of the coruses, number of active enrollment_id, which target variables are 1 in the training set, enrolled l
52	For each 14 days before the end date of the coruses, number of active enrollment_id, which target variables are 0 in the training set, enrolled l
53	For each 14 days before the end date of the coruses, number of active enrollment_id (in this case, days between last access and the end date o
54	For each 14 days before the end date of the coruses, number of active enrollment_id (in this case, days between last access and the end date o

Table 1: List of features generated by Song and Kohei.

ID	Model	Feature	5-CV	Pub
S1	xgb_rf_ko_new_feat	-	0.906721	0.90
S2	ko_v83	Feature RW + Feature SK + Feature JL + ko_feat3	0.906729	0.90
S3	bag_10_xgb_rf.xiv	Feature RW + Feature SK + Feature TN	0.905875	0.90
S4	xgb_rf.xvi	Feature RW + Feature TN + Feature JL	0.905543	0.90
S5	xgb_rf.xix	Feature RW + Feature SK + Feature TN + Feature MB	0.905356	-
S6	xgb_rf.xv	Feature RW + Feature TN	0.904935	0.90
S7	xgb_rw_sk	Feature RW + Feature SK	0.904914	-
S8	mjahrer.feat.tam+rw+sk+tam2+sk2+azure.dae+nn	Feature RW + Feature SK + Feature TN + Feature MJ	0.904235	-
S9	mjahrer.featmjahrer+Tam+RW+KoheiSong.nn	Feature RW + Feature SK + Feature TN + Feature MJ	0.903736	-
S10	mjahrer.featTam+RW+KoheiSong.dae+nn	Feature RW + Feature SK + Feature TN	0.903669	-
S11	pred_train_sk_feature_ffm_rw	-	0.903560	0.90
S12	mjahrer.featTam+RW+KoheiSong.nn	Feature RW + Feature SK + Feature TN	0.903428	-
S13	xg_400_4_0.05_feature_rw	Feature RW + Feature JL	0.903385	-
S14	kohei_song	Feature SK	0.902918	-
S15	xgb_cv_rw	Feature RW	0.902287	-
S16	ffm_cv_rw	Feature RW	0.901983	-
S17	mjahrer.dae+nn.RWfeat	Feature RW	0.901846	0.90
S18	mjahrer.featmjahrer+Tam+RW+KoheiSong.dae+krr	Feature RW + Feature SK + Feature TN + Feature MJ	0.901522	-
S19	xg_400_4_0.05_feature_sk	Feature SK	0.900906	-
S20	xgb_rf.xiv	Feature TN	0.899239	-
S21	xgb_rf.xiv	Feature TN	0.899167	-
S22	xgb_rf.xiv	Feature TN	0.898969	-
S23	xgb_rf.xiv	Feature TN	0.898890	-
S24	xgb.xiv	Feature TN	0.898749	-
S25	libfm_100_4_0.002_feature_rw_v2	Feature RW + Feature JL	0.898308	-
S26	xgl_500_0.5_10_10_feature_rw_v2	Feature RW + Feature JL	0.897968	-
S27	xgb_val.xiii	Feature TN	0.897912	-
S28	nn_20_16_0.01_feature_rw_v2	Feature RW + Feature JL	0.897143	-
S29	mb_nn_50_20_feat19	Feature RW + Feature MB	0.896748	-
S30	ffm_30_20_0.01_feature_rw_v2	Feature RW + Feature JL	0.896160	-
S31	xg_400_4_0.05_feature_tam	Feature TN	0.895754	-
S32	ConfigAMLKRR	Feature MJ	0.894524	-
S33	gbm.xiii	Feature TN	0.893507	-
S34	xg_600_4_0.05_feature10	Feature JL	0.892364	-
S35	xg_600_4_0.05_feature9	Feature JL	0.892253	-
S36	ConfigAML	Feature MJ	0.891217	-
S37	lr_0.01_feature_tam	Feature TN	0.890580	-
S38	ConfigAMLCUDAPreModel	Feature MJ	0.890565	-
S39	ffm_30_20_0.01_feature_tam	Feature TN	0.888418	-
S40	libfm_100_4_0.002_feature_tam	Feature TN	0.888381	-
S41	rf.xiii	Feature TN	0.887583	-
S42	et_1000_20_feature_tam	Feature TN	0.887768	-
S43	ffm_30_20_0.01_feature9	Feature JL	0.887116	-
S44	libfm_100_4_0.002_feature9	Feature JL	0.886866	-
S45	ConfigAML	Feature MJ	0.886705	-
S46	nn_20_16_0.01_feature9	Feature JL	0.886109	-
S47	xg_400_4_0.05_feature6	Feature JL	0.885184	-
S48	xg_400_4_0.05_feature3	Feature JL	0.885124	-
S49	ffm_20_20_0.01_feature6	Feature JL	0.885037	-
S50	ffm_20_20_0.01_feature3	Feature JL	0.884697	-
S51	xg_400_4_0.05_feature_mj	Feature MJ	0.882441	-
S52	et_1000_20_feature_rw_v2	Feature RW + Feature JL	0.881539	-
S53	libfm_100_8_0.01_feature6	Feature JL	0.880878	-
S54	libfm_100_8_0.01_feature3	Feature JL	0.880366	-
S55	nn_20_16_0.005_feature3	Feature JL	0.880219	-
S56	ConfigAMLKNN	Feature MJ	0.877652	-
S57	nn_20_16_0.005_feature5	Feature JL	0.876905	-
S58	lr_0.01_feature_mj	Feature MJ	0.821332	-
S59	lr_0.01_feature9	Feature JL	0.804150	-
S60	lr_0.01_feature3	Feature JL	0.802138	-
S61	lr_0.01_feature6	Feature JL	0.800225	-