Multi-Stage Ensemble and Feature Engineering for MOOC Dropout Prediction

Jeong-Yoon Lee
Conversion Logic
12300 Wilshire Blvd. Los Angeles, CA 90025, USA
jeong@conversionlogic.com

Andreas Toescher
Opera Solutions
Hauptplatz 12, 8580 Koeflach, Austria
andreas.toescher@commendo.at

Michael Jahrer
Opera Solutions
Hauptplatz 12, 8580 Koeflach, Austria
michael.jahrer@commendo.at

Kohei Ozaki
AIG Japan Holdings KK
Kamiyacho MT Bldg 18F, 4-3-20 Toranomon, Minato-ku, Tokyo 105-0001, Japan
ozaki.kohei@aig.co.jp

Mert Bay
Conversion Logic
12300 Wilshire Blvd. Los Angeles, CA 90025, USA
mert@conversionlogic.com

Peng Yan
NetEase Youdao
2nd Floor, Chuangye Building, Tsinghua Science Park, Beijing, 100084, China
yanpeng@rd.netease.com

Song Chen
American International Group, Inc. (AIG)
175 Water Street, New York, NY 10038, USA
song.chen@aig.com

Tam T. Nguyen
Institute for Infocomm Research, A*STAR
1 Fusionopolis Way, #21-01 Connexis (South Tower), Singapore 138632
nguyentt@i2r.a-star.edu.sg

Xiaocong Zhou
Tsinghua University
Haidian District, Beijing, 100084, China
infinitezxc@gmail.com

ABSTRACT

This paper describes the winning solution of KDD Cup 2015. The competition aims to predict dropouts in Massive Open Online Courses (MOOCs). Our approach begins with feature engineering to extract predictive features from activity logs of students and meta data. Then, we train sixty three individual classifiers with different subsets of features and seven algorithms. Lastly, we blend predictions of individual classifiers with the multi-stage ensemble framework. Our solution achieves AUC scores of 0.90918 and 0.90744 on the public and private leaderboards respectively.

Categories and Subject Descriptors

I.5.4 [**Pattern Recognition**]: Application

General Terms

Application

Keywords

KDD Cup, Feature Engineering, Ensemble Learning

1. INTRODUCTION

Our final solution is a joint work from 9 data scientists, distributed around the world. The pipeline from raw data to final solution is as follows:

- Hand crafted feature engineering (most of hard work)
- Automatic feature design (autoencoder)
- Individual models (gbm, nn, factor model,..)
- Stage-I ensemble (blends individual models)
- Stage-II ensemble (blends stage-I ensemble models)
- Stage-III ensemble (blends stage-II ensemble models)

2. DATASET

We got history from 200k enrollments, from 120k we know the labels. bla bla
bla bla

3. FEATURE ENGINEERING

All features.

3.1 Feature SK

Features generated by Song and Kohei can be classified as follows:

- Enrollment-based features (No.1-8)
- Username-based features (No.9-18)
- Username-based features for each courses (No.19-25)
- Features based on 10 days after the end date of course (No.26-35)
- Features based on 1 day after the end date of a course (No.36-45)
- Day-level features (No.46)
- Day-level features using target variables (No.47-58)

Full list of features generated by Song and Kohei are described in Table 1. (just listing them for now. TBD in detail).

3.2 Feature RW

Peng and Xiaocong feature.

3.3 Feature TN

Tam feature.

3.4 Feature MJ

Features generated by Michael Jahrer are in sparse format:

- uID(0-112447)
- cID(112448-112486)
- uIDcnt(112487-112487)
- eIDcnt(112488-112488)
- eID $\rightarrow$ sID(112489-112490)
- eID $\rightarrow$ evID(112491-112497)
- eID $\rightarrow$ oIDCnt(112498-139443)
- eID $\rightarrow$ tIDCnt(139444-139635)
- uID: floor(log(dateSpan^2+1))(139636-140635)
- uID $\rightarrow$ log(time diff to obj start+1)(140636-140636)
- eID $\rightarrow$ dateVec diff stats(140637-140649)

3.5 Feature MB

Mert feature.

3.6 Feature JL

Jeong feature.

4. INDIVIDUAL MODELS

something like this.

4.1 Learning Algorithms

- Logistic Regression (LR)
- Kernel Ridge Regression (KRR)
- Factorization Machine/Field-aware Factorization Machine (FM/FFM)
- Neural Networks (NN)
- Extreme Trees (ET)
- Gradient Boosting Decision Trees (GBDT)

4.2 Individual Models

- Model 1: RandomForest(R). Dataset: X
- Model 2: Logistic Regression(scikit). Dataset: Log(X+1)
- Model 3: Extra Trees Classifier(scikit). Dataset: Log(X+1) (but could be raw)
- Model 4: KNeighborsClassifier(scikit). Dataset: Scale(Log(X+1))
- Model 5: libfm. Dataset: Sparse(X). Each feature value is a unique level.

5. MULTI-STAGE ENSEMBLE

Stratified 5-fold cross validation (CV). We used xgboost [1], neural nets and linear regression for stage-II ensembling.

5.1 Stage-I Ensemble

We trained 20 stage-I ensemble classifiers with different subsets of CV predictions of 63 individual classifiers.

5.2 Stage-II Ensemble

We trained 2 stage-II ensemble classifiers with different subsets of CV predictions of 20 stage-I ensemble classifiers.

5.3 Stage-III Ensemble

We trained a stage-III ensemble classifier with CV predictions of 5 classifiers: 2 stage-II ensemble, 2 stage-I ensemble, and 1 individual classifiers.

id	name	type	5CV
1	trn.final.90788	single?	0.907878
2	esb58v5+magic.dae+nn.validCV.0.907567	Stage-I	0.907567
3	et_esb58v5_rank.val.0.906207	Stage-I	0.906207
4	lr_forward_0.01_esb.esb15v3.val.yht	Stage-I	0.907968
5	xgb_rf.ko_new_feat.txt.valCV.0.906721	single?	0.906721

A linear combination of the 5 models from table 5.3 results in train AUC=0.908072 and accuracy=0.887334. Which leads to 0.90910 public leaderboard score. By adding 39 courseID correction factors train AUC=0.908194 and public score improved to 0.90918.

6. CONCLUSIONS

Our final AUC score of 0.90918 results from a complex pipeline from raw data to final score. Every part of that pipe needs to be (sub-)optimal implemented by our team to get the best score at the end. The first part "feature design" is the most important one and needs expertise, experience and of course a bit luck to capture all signals in the data.

7. ACKNOWLEDGEMENTS

Thanks to dropbox, github and skype to enable easy communication around the globe.

8. ADDITIONAL AUTHORS

how we can get rid of this ? this section appears with more than 6 authors.

9. REFERENCES

[1] xgboost. https://github.com/dmlc/xgboost.

No.	Description
1	Course_id encoded by 1-of-N coding
2	Number of requests by an enrollment_id
3	Number of unique object by an enrollment_id
4	Number of unique problem object of event by an enrollment_id
5	Number of active days by an enrollment_id
6	Number of active hours by an enrollment_id
7	Time of first access in hours by an enrollment_id
8	Time of last access in hours by an enrollment_id
9	Number of enrollments by an username
10	Number of requests by an username
11	Number of unique objects by an username
12	Number of unique problem object of event by an username
13	Number of active days by an username
14	Number of active hours by an username
15	Time of first access in hours by an username
16	Time of last access in hours by an username
17	Time of first problem access in hours by an username
18	Time of last problem access in hours by an username
19	For each course, number of requests by an username
20	For each course, number of unique object by an username
21	For each course, number of unique problem object by an username
22	For each course, number of active days by an username
23	For each course, number of active hours by an username
24	For each course, time of first access in hours
25	For each course, time of last access in hours
26	Number of enrollment_ids during 10 days after the end date of course by an username
27	For each course, number of access logs during 10 days after the end date of course by an username
28	For each course, number of unique objects during 10 days after the end date of course by an username
29	For each course, number of unique problem objects during 10 days after the end date of course by an username
30	For each course, number of active hours during 10 days after the end date of course by an username
31	For each course, difference between first and last access during 10 days after the end date of course by an username
32	For each course, time of first access in hours during 10 days after the end date of course by an username
33	For each course, time of last access in hours during 10 days after the end date of course by an username
34	For each course, time of first access to an problem object in hours during 10 days after the end date of course by an username
35	For each course, time of last access to an problem object in hours during 10 days after the end date of course by an username
36	Number of enrollment_ids during 1 day after the end date of course by an username
37	For each course, number of access logs during 1 day after the end date of course by an username
38	For each course, number of unique objects during 1 day after the end date of course by an username
39	For each course, number of unique problem objects during 1 day after the end date of course by an username
40	For each course, number of active hours during 1 day after the end date of course by an username
41	For each course, difference between first and last access during 1 day after the end date of course by an username
42	For each course, time of first access in hours during 1 day after the end date of course by an username
43	For each course, time of last access in hours during 1 day after the end date of course by an username
44	For each course, time of first access to an problem object in hours during 1 day after the end date of course by an username
45	For each course, time of last access to an problem object in hours during 1 day after the end date of course by an username
46	For each days of the course, which date is provided in date.csv, number of unique active courses by an username
47	For each 10 days after the end date of the course, number of active enrollment_id, which target variables are 1 in the training set, enrolle
48	For each 10 days after the end date of the course, number of active enrollment_id, which target variables are 0 in the training set, enrolle
49	For each 10 days after the end date of the course, number of active enrollment_id (in this case, days between last access and the end date
50	For each 10 days after the end date of the course, number of active enrollment_id (in this case, days between last access and the end date
51	For each 14 days before the end date of the coruses, number of active enrollment_id, which target variables are 1 in the training set, enro
52	For each 14 days before the end date of the coruses, number of active enrollment_id, which target variables are 0 in the training set, enro
53	For each 14 days before the end date of the coruses, number of active enrollment_id (in this case, days between last access and the end da
54	For each 14 days before the end date of the coruses, number of active enrollment_id (in this case, days between last access and the end da

Table 1: List of features generated by Song and Kohei.